

Jacopo Zacchigna

MSc Student · Mechanistic Interpretability of LLMs · Trieste, IT

✉ jaczac2002@gmail.com | 🌐 jac-zac.github.io | 🗣️ [Jac-Zac](#) | 📄 [jacopo-zacchigna](#) | 🎓 [Google Scholar](#)

Research Experience

Supervised Program for Alignment Research (SPAR), *Research Fellow* Feb 2026 – May 2026

Research fellowship under *Gabriele Sarti* studying persona representations in LLMs ([implicit-personalization](#)).

- Released [persona vectors](#) for Gemma-2-27B and Llama-3.1-405B to probe and steer user representations.
- Developed [persona-data](#) with loaders and prompt utilities for the [SynthPersona](#) dataset (1k personas).
- Built an interactive [web UI](#) for real-time extraction, analysis, and probing of implicit persona representations.

Mechanistic Interpretability of Mixture-of-Experts LLMs, *Independent research* Jan 2026 – present

- Adapting HeadPursuit / SOMP to identify expert specialization (e.g., toxicity) in OLMoE and GPT-OSS.
- Investigating expert editing and steering interventions versus knockout-style ablations.

Selected Projects

Causal Multi-Head Self-Attention Kernels, [code](#) Jan – Feb 2026

- Implemented causal MHSA in CUDA, OpenMP, and SIMD on Leonardo (Cineca), iterating six CUDA kernels up to an online-softmax (FlashAttention-style) version.
- Scaled the OpenMP kernel to 128 cores with 94× speedup over its single-threaded baseline (73% parallel efficiency), 19× faster than naive PyTorch.

BayesianFlow Uncertainty Estimation in Generative Models, [code](#) Jun – Jul 2025

- Extended BayesDiff to flow matching, achieving 5× **faster** generation than DDIM at comparable quality.
- Integrated Last-Layer Laplace Approximation for pixel-wise uncertainty (*interpretable* confidence maps).

Publications & Awards

Zacchigna, J. & Liu, W. (co-first authors), et al. *Application of Computer Vision to the Automated Extraction of Metadata from Natural History Specimen Labels: A Case Study on Herbarium Specimens*. [Plants \(MDPI\)](#) 2026

[Aletheia's Quest](#), LLM lie-detection competition: 3rd place, black-box track (Team JSR) 2026

Education

University of Trieste, MSc, Data Science & Artificial Intelligence – Trieste, Italy 2024 – 2027
Relevant coursework: Theory of Deep Neural Networks, Explainable AI, Probabilistic ML, Advanced Deep Learning.

University of Trieste, BSc, AI & Data Analytics – Trieste, Italy 2021 – 2024

Experience

Obloo, AI Consultant – Trieste, Italy May 2024 – Jan 2025

- Designed and deployed retrieval-augmented generation (RAG) systems over the company data lake.
- Mentored 5 interns to independent project delivery; led ML prototypes adopted in two company-wide products.

Skills

Technical: Python, C/C++, CUDA, PyTorch, nnsight/nninterp, Hugging Face, NumPy, R, Docker, Linux, HPC

Languages: Italian (native), English (C1)